

ScienceWorld: environment inspection report

Wave-0 M5, CA-Benchmark project

2026-07-23

Scope. Hands-on inspection of the ScienceWorld text environment (rung 4 of the credit-assignment benchmark): what the tasks are, how they are scored and evaluated, how many trials exist, rollout speed on our node, the sandbox/action surface, a full worked trajectory, expert-data availability, published baselines, and the challenges that bear on training. All local numbers are reproduced by `scripts/scienceworld/audit_structure.py`, `analyze_audit.py`, and `bench_and_probe.py`; raw assets are in `assets/`. Package: `scienceworld 1.2.3` (pip) on OpenJDK 17, one CPU worker.

1 What the environment is (the sandbox)

ScienceWorld (Wang et al., EMNLP 2022; arXiv 2203.07540) is a **text-based interactive simulator of elementary-school science**. The agent navigates a map of **10 locations** centered on a house (hallway, kitchen, bathroom, workshop, art studio, greenhouse, bedroom, living room, foundry, outside/ garden), populated with **up to 200 object types**, carrying out science procedures. It is **partially observable**: `look around` reveals only the current room; the underlying engine tracks a full thermodynamic / chemical / biological state (temperatures, states of matter, electrical circuits, plant and animal life-cycles) that the agent must actively probe with instruments.

Engine & interface. The simulator core is **~40k lines of Scala** (compiled to a JAR, run on Java 1.8+) exposed to Python via `py4j` — a JVM must be present (we installed OpenJDK 17). Verified from the source: each `ScienceWorldEnv` launches **its own dedicated JVM process** (plus a Python callback server); there is no documented thread-safety mechanism, so **parallelism is process-level** (the paper’s RL runs used 8 parallel env instances). For us this means the environment is CPU/JVM-bound and single-threaded per instance; worker throughput scales with the number of env processes we co-locate with the trainer, not with threads.

Action space (25 templates). Actions are *template + object(s)* combinations:

activate/deactivate OBJ	open/close OBJ	focus on OBJ
go OBJ	teleport OBJ	look around / look at / look in OBJ
pick up / put down OBJ	move OBJ to OBJ	pour OBJ in OBJ
dunk OBJ in OBJ	mix OBJ	use OBJ on OBJ
read OBJ	eat OBJ	flush OBJ
wait / wait1	inventory	task

The **combinatorial action space** is the headline difficulty: the paper estimates **~200,000 naively-possible action-object combinations per step**; in our own probe a single `boil` start-state already exposes **115 valid combinations**, growing as the agent uncovers objects. A free-text agent (our ReAct policy) is *not* handed this list — it must generate a valid `Thought:/Action:` string, and malformed or inapplicable actions get a soft error observation (e.g. "You can't move something into itself."). Critically, the paper reports that all but two of its baselines relied on a **valid-action-detection aid** at test time to shrink this search — a crutch our generative agent does not use, which makes valid-action-rate a first-class diagnostic. **Simplifications** can be toggled; the audit used the standard `easy` set (`teleportAction`, `openDoors`, `selfWateringFlowerPots`, `noElectricalAction`) — these must be frozen identically across all arms and recorded, since they change both horizon and answerability.

2 The tasks and how they are scored (the evaluation)

There are **30 tasks** (10 topic categories) spanning thermodynamics, chemistry/mixing, electricity, classification, biology (life stages, lifespans, plant growth, genetics), and measurement; we group them into **10 families** (`assets/families.csv`). The paper reports **7,200 total variations across all 30 tasks** — different target substances, objects, and layouts. Each task ships fixed **train / dev / test** variation splits ($\approx 50/25/25\%$), so “unseen variations” is a built-in generalization split — no custom splitting needed. The original protocol resets an episode at success, failure, or a **100-step cap**; we replace that single global cap with per-family caps (below), since horizons range $\sim 40\times$.

Scoring is a dense, cumulative 0–100 subgoal score. The per-step **reward** returned by the API is the score *delta*. From replaying 90 gold trajectories (3 train variations $\times$ 30 tasks):

- All 90 gold paths reach exactly **score 100 with done=True**; no negative deltas on-path.
- Subgoal firing density: median **0.27** score-increments per step (range 0.05–0.62) — dense relative to a terminal-only signal, but far from every-step.

- **Magnitudes are heterogeneous and task-specific: 42 distinct values from 1 to 72.** The oracle is *not* a uniform unit-per-subgoal reward. Any process-reward (PRM) arm in the side experiments must match this granularity or renormalize explicitly and say so.
- **Off-gold failure is catastrophic:** a wrong *focus* on (focusing on the wrong object) returns **reward −100 and immediately terminates the episode.** *focus* is thus both a hard subgoal gate and an instant-death / reward-hacking surface for exploration.

Our reward convention (proposed, confirm at W1 manifest). Core-matrix arms train on *terminal success only*: success $\Leftrightarrow$ final score = 100; a −100 focus-death maps to terminal reward 0 and episode end (mirroring ALFWorld failure, keeping terminal reward in {0,1}). The dense 0–100 oracle is reserved for the PRM-ladder / redistribution side experiments, where the −100 event is exactly the kind of asymmetry the redistribution-vs-separation question probes.

Oracle isolation (checked). Observation text on all 90 gold paths contained zero **score/reward/subgoal/points** strings; the score lives only in the API **info** dict, and the **task** command exposes only the task description (already in the prompt). So a transcript-only agent is blind to the oracle — a prerequisite for the “rubric PRM must not see oracle state” guarantee. (One residual: audit our own env *wrapper* once written.)

3 How many trials, and rollout speed

Trials. 30 tasks with widely varying variation counts — from 4 train variations (**identify-life-stages-2**) to **692** (**inclined-plane-friction-named-surfaces**). The **find-*** and **test-conductivity** families have 150–450 train variations each. The paper’s **7,200 total variations** (train+dev+test across all tasks) means that, unlike a fixed-size QA set, ScienceWorld supplies a large pool of fresh task instances — the training-time question is the *sampling* policy (uniform-per-task vs uniform-per-variation), which we must fix before Wave 1 because the 4-vs-692 imbalance changes what one “epoch” means and which tasks dominate the gradient.

Rollout speed (single JVM env, our node; **assets/speed_bench.json**):

Operation	measured
<code>load(task,var) + reset</code>	0.22 s/episode (0.15–0.36)
Gold-action stepping	13.8 steps/s (73 ms/step)
Random stepping, step only	15.3 steps/s (65 ms/step)
Random stepping incl. <code>get_valid_actions</code> each step	9.0 steps/s (112 ms/step)

Reading for training budget. At $\approx 65\text{--}75$ ms/env-step, one env instance yields ~ 14 steps/s. The horizon is the multiplier that matters: a 30-step task costs ~ 2 s of pure env time, but a 200-step **inclined-plane** episode costs ~ 14 s of *CPU env time alone*, before any LLM generation. This is why ScienceWorld is the most expensive rung per environment-episode and why long-horizon families dominate wall-clock. Querying the valid-action set every step (which we do *not* need for a free-text agent) nearly halves throughput — keep it out of the hot loop. Env parallelism is process-level (one JVM each), so worker-level throughput scales with the number of env workers we run alongside the trainer.

4 A full trajectory (gold/expert)

Expert data is available for free: the engine generates a **gold action sequence** for any task/variation (`generateGoldPath=True  $\rightarrow$  get_gold_action_sequence()`). This is our rung-4 SFT/ π_{base} source (rendered into the ReAct grammar, with the score channel stripped — oracle isolation). No external expert dataset is needed. Below is a complete gold trajectory for a short task; green rows mark subgoal firings (with the reward increment), showing how the dense signal is distributed along the path.

Task find-animal (variation 0), 10 gold steps, final score 100.

“Your task is to find a(n) animal. First, focus on the thing. Then, move it to the red box in the kitchen.”

<i>t</i>	action	observation	Δscore
0	open door to greenhouse	The door is already open.	+8
1	go to greenhouse	You move to the greenhouse.	+9
2	open door to outside	The door is already open.	
3	go to outside	You move to the outside.	
4	look around	...an axe, a blue jay egg, ...(many objects)	
5	focus on blue jay egg	You focus on the blue jay egg.	+50
6	pick up blue jay egg	You move the blue jay to the inventory.	+8
7	open door to kitchen	The door is already open.	
8	go to kitchen	You move to the kitchen.	+8
9	move blue jay egg to red box	You move the blue jay to the red box.	+17

A second, richer trajectory (**change-the-state-of-matter-of** water, 22 steps, in **assets/full_traj_*.json**) shows the harder physics: the agent picks up a thermometer, puts water in a metal pot, **focus on substance in metal**

pot (+66, the big subgoal gate), moves it to the freezer (+2), then **waits** while probing temperature (10°C → 3°C) until the water freezes and the final **wait** fires +29 to reach 100. It illustrates three training-relevant traits at once: **partial observability** (state read only via the thermometer), **delayed physical dynamics** (must **wait** for freezing — reward arrives many steps after the causal action), and a **dominant single subgoal** (focus worth 66/100), which shapes what credit assignment has to discover.

5 Published baselines (original paper + follow-ups)

Original paper (Table 2; test variations; 0–1 scale). The headline finding is that *elementary science is hard for text agents and larger models are not better*: a 1.5M-parameter RL agent trained interactively for 100k steps beats an 11B model trained statically for science QA.

Agent	Params	Avg score (0–1)
Random-valid	—	0.03
DRRN (best baseline)	1.5M	0.17
KG-A2C	5.5M	0.11
CALM (GPT-2)	131M	0.05
Behavior Cloning (T5-11B)	11B	0.08
Text Decision Transformer (T5-11B)	11B	0.08

Best agent 0.17/1.0 — comparable to text agents on medium-difficulty interactive-fiction games. The paper notes all but two baselines leaned on the valid-action aid at test time. Source: arXiv 2203.07540 Table 2.

Follow-ups (0–100 scale; note the differing eval sets — not directly comparable across rows). Modern LLM agents changed the picture dramatically:

Method	Base model	Eval set	Avg score
DRRN / KG-A2C / CALM	(paper baselines)	270 test vars	15.6 / 11.4 / 5.1
SayCan	GPT-4	270 test vars	33.8
ReAct	GPT-4	270 test vars	36.4
Reflexion	GPT-4	270 test vars	45.3
SwiftSage	T5-large (770M) + GPT-4	270 test vars	84.7
PPO	Llama-2-7B-Chat	194 seen / 241 unseen	59.4 / 51.7
SFT	Llama-2-7B-Chat	194 seen / 241 unseen	67.4 / 53.0
ETO	Llama-2-7B-Chat	194 seen / 241 unseen	73.8 / 65.0
GPT-4 (ETO paper)	GPT-4	194 seen / 241 unseen	64.8 / 64.4
BC-large	Llama-2-7B-Chat	AgentGym split	74.5
AgentEvol	Llama-2-7B-Chat	AgentGym split	38.0
GPT-4-Turbo (AgentGym)	GPT-4-Turbo	AgentGym split	14.4

Sources: SwiftSage arXiv 2305.17390 (first 10 test variations/task = 270); ETO arXiv 2403.02502 / ACL 2024 (avg reward, seen/unseen); AgentGym arXiv 2406.04151. **Reading for us:** (1) a 7B open model with the right training (ETO 65–74, SFT 53–67) is a realistic target band for our Qwen3-4B rung-4 numbers — meaning 4B *can* get off the ground here, unlike the ASearcher worry. (2) SFT/BC is a *very* strong baseline on this env (BC-large 74.5, SFT-seen 67.4), so our SFT π_{base} floor rows matter enormously — an RL arm must clear a high bar, and AgentEvol *losing* to BC-large is a cautionary precedent. (3) The GiGPO/verl-agent and RAGEN lineages do **not** report ScienceWorld, so our rung-4 numbers are largely new territory for those exact methods — good for novelty, but no in-method reference point exists. **Caveat:** every follow-up uses a different eval subset (270 vs 194/241 vs AgentGym), so cross-row comparison is illustrative only; our SOTA table (§10 of the overview) will carry base-model + split + metric columns to avoid the mismatched-settings trap.

6 Challenges (why this rung matters for credit assignment)


1. **Long, variable horizons (5–200+ steps).** Sampled gold-path lengths span $\sim 40\times$ within one rung (see per-family table), giving RQ2 within-rung horizon strata for free — but also making this the hardest rung for flat trajectory-level credit, which is exactly the hypothesis under test.
2. **Combinatorial action space** (>100 valid combos/step) with soft failure on invalid actions — a free-text agent must *generate* legal actions, so a valid-action-rate metric is a leading diagnostic (as it was in the ALFWorld pilot).
3. **Partial observability + delayed dynamics.** Rewarding actions (focus, heat, wait) often pay off many steps later, and world state must be actively measured. This is the regime where a turn-level critic or anchor-state advantage should, in principle, beat token-broadcast GRPO.
4. **Catastrophic focus death (–100, instant end).** A per-turn hacking/instant-fail surface; needs a focus-death-rate counter and a decided terminal mapping.

5. **Heterogeneous dense reward (magnitudes 1–72, one subgoal often dominates).** Makes the naive “sum dense + terminal” channel spiky — the motivation for the redistribution-vs-separation side experiment.
6. **Context blow-up on long families.** Proposed $2\times$ -gold turn caps reach 260–404 for genetics / matter-state / inclined-plane; full-history transcripts will not fit at 4B, so oldest-observation truncation is load-bearing for 4 of 10 families. Final caps await strong-model turns-to-success calibration.
7. **Task-sampling / variation imbalance (4 vs 692 variations).** Must be fixed a priori so “epoch” and the train/unseen generalization gap are well-defined.

Per-family horizons and proposed turn caps

family	sampled gold-len range	proposed cap ($2\times\max$)
lifespan	6–9	18
find-thing	5–16	32
chemistry-mix	11–30	60
life-stages	10–30	60
electricity	9–45	90
plant-growth	66–85	170
thermal-measure	18–94	188
genetics	123–130	260
matter-state	20–143	286
inclined-plane	53–202	404

Caps are provisional ($2\times$ the longest sampled gold path); they will be finalized from strong-model (gpt-oss-120b) rollout turns-to-success (95th percentile) rather than the gold heuristic alone, since gold paths contain filler an agent may skip or exceed. Annotated HTML trajectories in `assets/html/`.



Gold-path length (steps) by task family, 3 sampled train variations each — the $\sim 40\times$ within-rung horizon spread that gives RQ2 its horizon strata.