

Throughput Is Not Learning: Final Report on the Async-Rollout 8B Campaign

2× wall-clock validated; per-sample parity at best; transfer deficits traced to answer-coverage collapse

CA-benchmark rung-3 / 8B thread

2026-08-14

Abstract

We complete the four-arm 8B campaign crossing two rollout engines (old lockstep vs. a trajectory-concurrent async collector) with two credit-assignment methods (token-GRPO, turn-PPO) on ASearcher search-RL, 75 steps each, plus a 7,152-question transfer evaluation under three scorers. The async engine’s systems claims hold: **2.01×** end-to-end at matched config, 3–4× on the rollout phase, zombie generation 49.6%→0. Per training *step* it learns ~2× faster early (val 0.43 vs. 0.23 at step 25) — but sample-matched curves show this is a *data-rate* effect: at equal cumulative released trajectories the old engine matches or beats it, and on transfer the old arms win everywhere (macro judge: old turn-PPO **0.550**, old GRPO 0.535, async GRPO 0.506, async turn-PPO 0.445). Decomposition localizes the async deficits almost entirely to *answer coverage*: conditional on answering, async turn-PPO is the most accurate arm in the campaign (judge 0.623), but it emits no answer on 30.2% of questions (old: 0.3%), terminating via context exhaustion after long searches; async GRPO shows the same drift at 1/3 strength. Stale-group drops were zero in both engines — staleness censoring is *not* a mechanism. We flag the critic × partial-rollout-resume interaction as the prime suspect for the late-training regression (val 0.486→0.404 over steps 50–75) and recommend a staleness correction plus answer-commitment shaping before the async engine becomes the training default.

Setup

Qwen3-8B-Base, unfiltered ASearcher Base-35k, 32-turn horizon, 16k prompt / 1024 response, top-5 wiki-18/e5 retrieval, 75 steps, both engines with identical partial-rollout config (cycle 8 / max_age 4). Async arms (**async75kv60**): one coroutine per trajectory over vLLM V1 AsyncLLM engines, uid-atomic group release, gpu_util 0.60, sticky routing. Transfer evaluation: greedy **val_only** on the 10-benchmark suite; 7 wiki-answerable benchmarks (6,115 q) reported separately from 3 live-web benchmarks (1,027 q, corpus gap); scorers: strict EM (the training reward), substring EM, gpt-oss-120b judge (frozen MBE prompt). All four arms evaluated on the same frozen old-engine harness. Caveats: async-arm evals ran on A100 (old on H100), cross-arch greedy drift ≤0.02 — deficits are 3–8× that; judge ≈ EM+0.13, one-directional.

1 Engine mechanisms (matched profile, corrected)

metric	sync	async	ratio
phase wall-clock (3 steps + val + ckpt)	8,953 s	4,454 s	2.01×
rollout generation, 3-step total	4,890 s	1,228 s	~3–4×
zombie generated tokens	49.6%	0	—
released trajectories	2,300	2,695	1.17×
val@3 (greedy, 512 q)	0.264	0.244	within ±0.02 band

Table 1: Matched H100 A/B (same config/seed/data). Steady state @0.65 util: 6,473 gen tok/s, 174 traj/min. B200 measured 0.88× H100 on rollout (TRITON_ATTN decode 1.6× slower/token) — H100 remains the rollout node.

Correction: **max_age** stale-group drops were **zero** in both engines across all 75 steps of all four arms — stale censoring explains nothing. The operative data-side differences: (1) **release volume**: async releases 1,200–1,700 traj/step vs. 690–1,150 (cumulative 92–94k vs. 55–69k), because the old engine’s fixed batch geometry drops

surplus completed groups; (2) **completion mix**: continuous slot backfill keeps long trajectories that lockstep cycles would truncate or never collect; (3) **fresher resumes**: paused trajectories resume immediately rather than at the next 8-turn cycle; (4) fp16 engine numerics (token ids byte-exact by construction).

2 Per-step vs. per-sample learning




Figure 1: Train success rate (EMA; faint = raw) and greedy val EM by *step*. Async arms are $\sim 2\times$ ahead at step 25. Dips near step 26 are pod-reclaim/OOM resume replays, not learning events.




Figure 2: The same val points re-indexed by cumulative released trajectories. The async per-step advantage is a data-rate effect: old GRPO reaches 0.520 by ~ 38 k trajectories; async GRPO needs ~ 92 k for 0.529. Async turn-PPO peaks at 0.486 (61k) then *falls* to 0.404 (94k).

3 Transfer: the ordering reverses

4 Failure analysis: answer coverage, not accuracy

Paired per-question analysis (async vs. old turn-PPO, 6,116 common questions): 1,028 lost, 318 gained; **720 of the 1,028 losses are no-answer terminations**. No-answer trajectories average 17.2 turns with **0%** at the 32-turn cap: they end by *context exhaustion* (`ctx_terminate`) — the policy searches until the 16k window fills without committing to an answer. The drift is training-time: async turn-PPO’s own val fell 0.486 $\rightarrow$ 0.404 (steps

arm	engine	val@0	val@25	val@50	val@75
old GRPO	lockstep	0.195	0.227	0.520	0.520
old turn-PPO	lockstep	0.195	0.234	0.479	0.561
async GRPO	async	0.203	0.434	0.418	0.529
async turn-PPO	async	0.199	0.322	0.486	0.404

Table 2: Greedy strict-EM validation (512 q, ± 0.02 reproducibility).




Figure 3: Wiki-answerable macro under three scorers, plus live-web judge. Old engine wins every cell; the async GRPO train-val advantage (0.529 vs. 0.520) reverses on transfer (0.506 vs. 0.535).

50–75) while train success kept rising. Prime suspect: the **critic** $\times$ **partial-rollout-resume interaction** — resumed trajectories carry returns realized under older policies; a turn-level critic bootstraps across the resume seam, and async’s higher resume volume makes the value of “continue searching” systematically stale-optimistic. GRPO (no critic) shows the same direction at $\sim 1/3$ strength. Candidate fixes: behavior-IS staleness correction (parked item R2), answer-commitment shaping, resume-age cap for critic arms.

5 Credit-assignment sweep (old engine, 5 estimators)

HCAPO: incipient per-turn length runaway. HCAPO is the only arm that never compresses per-turn responses (median 270–400 tok/turn throughout, vs. ~ 30 for GRPO/GiGPO and ~ 100 –120 for the PPOs; p90 pinned at the 1,024 response cap for the entire run). Over steps 62–75 the median climbs $\sim 280 \rightarrow 390$ with spikes to 716 (step 72) and **1,024** — **the cap** — **at step 75**, while the valid-action ratio sags $0.84 \rightarrow 0.79$ (cap truncation). This is mechanism-consistent: $\gamma=0.95$ discounts per *turn*, so content migrates into fewer, longer turns, and nothing in the hindsight objective penalizes within-turn tokens. Trajectory length is *not* running away (turns fell $4.8 \rightarrow 3.4$) and reward was unaffected through step 75, but longer training would likely saturate the cap; follow-ups should add a per-token length penalty or a per-turn cap margin.

Verdict

Adopt the async collector for what it is: a **2 $\times$ wall-clock engine** with clean semantics (atomic groups, zero zombie generation, val parity at profile scale). Do not yet adopt it as the training default: its data distribution (more, longer, fresher-resumed trajectories) pushes policies toward non-committal search, mildly for GRPO and severely for critic-based turn-PPO. Fix order: (1) staleness/IS correction at the resume seam, (2) answer-commitment shaping or no-answer penalty, (3) re-run the turn-PPO arm; the sample-matched curves here are the baseline any fix must beat.

arm	judge (all)	judge (answered)	no-answer	mean turns	median
old GRPO	0.535	0.571	6.2%	17.5	18
async GRPO	0.506	0.573	11.7%	18.7	19
old turn-PPO	0.551	0.552	0.3%	7.8	3
async turn-PPO	0.435	0.623	30.2%	8.5	3

Table 3: Wiki-answerable decomposition. Conditional accuracy is preserved or improved under the async engine; coverage is what collapses.




Figure 4: Left: all-question vs. answered-only judge score. Right: async turn-PPO no-answer concentrates on multihop benchmarks (Musique 43%, 2Wiki 42%, Hotpot 31%).




Figure 5: Turn distributions (wiki-answerable). Both turn-PPOs are bimodal (52–54% answer within 3 turns); async shifts the search mode right. Async GRPO inflates the 20–23 bucket to 35% (old: 10%). No arm hits the 32-turn cap.

estimator	mechanism	val@0	val@25	val@50	val@75	turns@75
turn-PPO (b0)	turn-level GAE + critic	0.195	0.234	0.479	0.561	7.5
GiGPO	anchor-state + episode groups	0.195	0.064	0.438	0.535	16.1
HCAPO	hindsight answer-cond., $\gamma=0.95$	0.195	0.297	0.449	0.529	3.7
GRPO	token-level group-relative	0.195	0.227	0.520	0.520	17.0
token-PPO	stock token GAE + critic	0.195	0.312	0.447	0.426	8.4

Table 4: Estimator comparison on the lockstep engine (**fast75**), all arms complete. token-PPO / HCAPO / GiGPO died with their pods at steps 42/44/46, were resumed 2026-08-14 from step 25 on identical code (restart re-vals reproduced 0.312/0.297/0.064 exactly), and finished 2026-08-15/16.




Figure 6: Five-estimator training curves, all complete. turn-PPO wins (0.561); GiGPO (0.535) and HCAPO (0.529) edge past GRPO (0.520); token-PPO trails (0.426). **GiGPO’s step-25 val collapse (0.064) fully recovered** — its train success was normal (~ 0.45) throughout; the dip was a transient greedy-decode pathology (17.8 val turns at step 25), not estimator failure. **token-PPO regressed $0.447 \rightarrow 0.426$ over steps 50–75** — both late regressions in the campaign (this and async turn-PPO, §5) are critic-based arms. HCAPO is the turn-efficiency standout: 0.529 at 3.7 turns/question, half of turn-PPO’s 7.5; its $\gamma=0.95$ is a deliberate protocol deviation (paper value). val@25 was weakly predictive of val@75: the step-25 leader finished last, the step-25 collapse finished second.