


ALFWorld (Qwen3-1.7B, ReAct transcript) — training

ALFWorld (Qwen3-1.7B, ReAct transcript) — final eval



method	unseen success	eval step
token-PPO	0.761	150
turn-PPO	0.963	150
GRPO	0.881	265
GiGPO	0.955	200
HCAPO	0.784	145
HCAPO (adapted)	0.791	145
pre-RL init	0.440	0